

# HTTP/3


What were we thinking?

# Reminder: HTTP/1.1



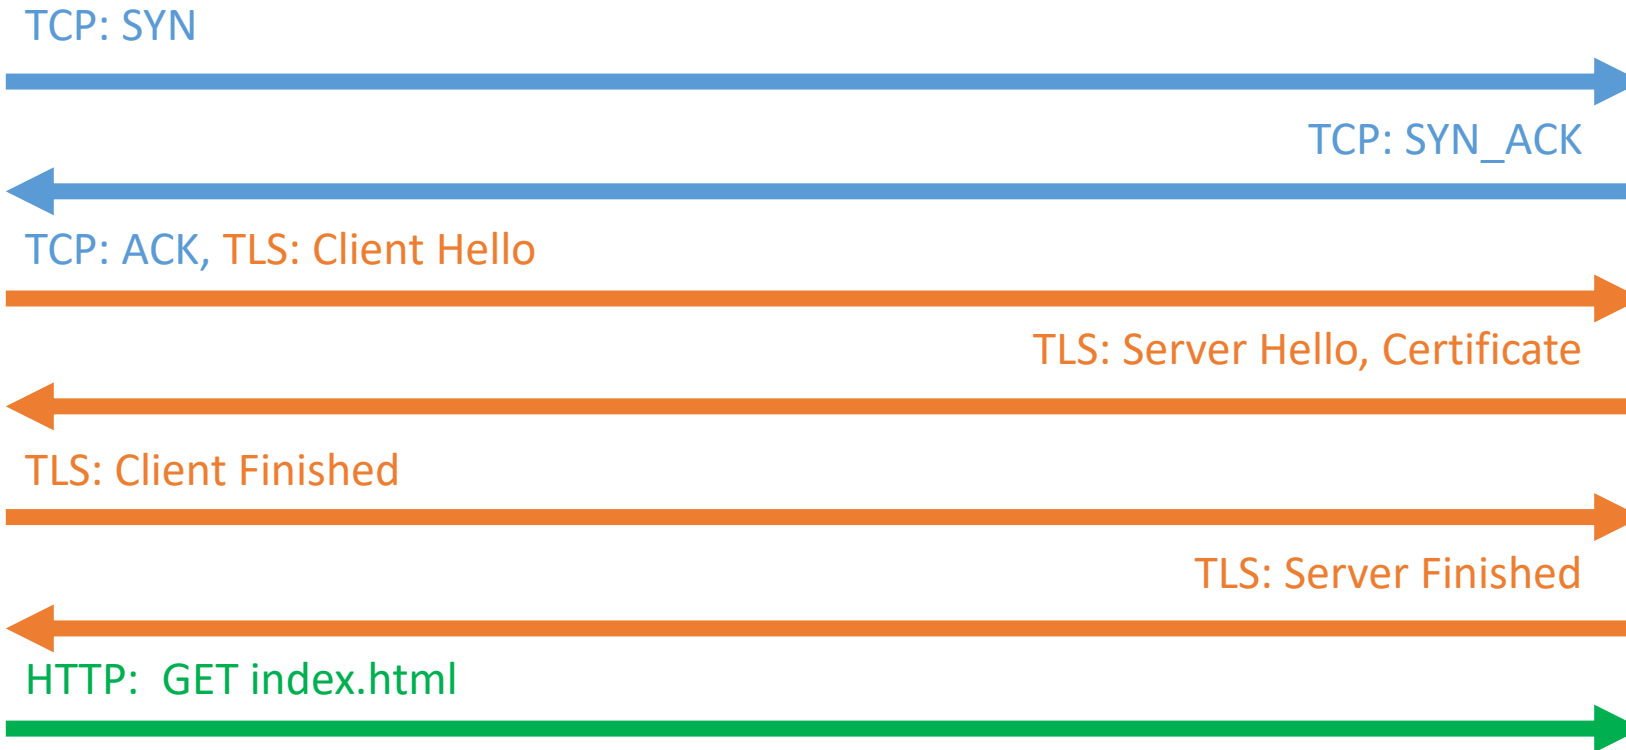
- One request/response at a time
- Multiple connections to make multiple requests
- Every connection pays connection setup costs

# HTTP/2



- Many requests and responses on a single TCP connection
  - Pay setup costs once
  - Congestion controller can manage connection better

# Connection Setup

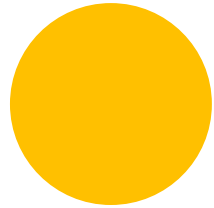



•  
•  
•

TCP Fast Open

TLS False Start or  
TLS 1.3

TLS 1.3 Early Data



So we're done, right?

...once we can deploy all those things, anyway....










# Deployment is a big problem

- Middleboxes inspect and tamper with non-zero probability on any port, any portion of the message
  - Makes deploying new protocols on the Internet very hard
- Some examples
  - Windows has tried to turn TCP Fast Open on by default in at least the last four releases, keeps having to revert it
  - Each browser has a built-in list of sites not to try TLS False Start with
  - So many devices break with TLS 1.3 that it lies and claims to be TLS 1.2 for the benefit of older devices
    - The actual version number is hidden in a new extension
- Middleboxes constrain innovation on the Internet

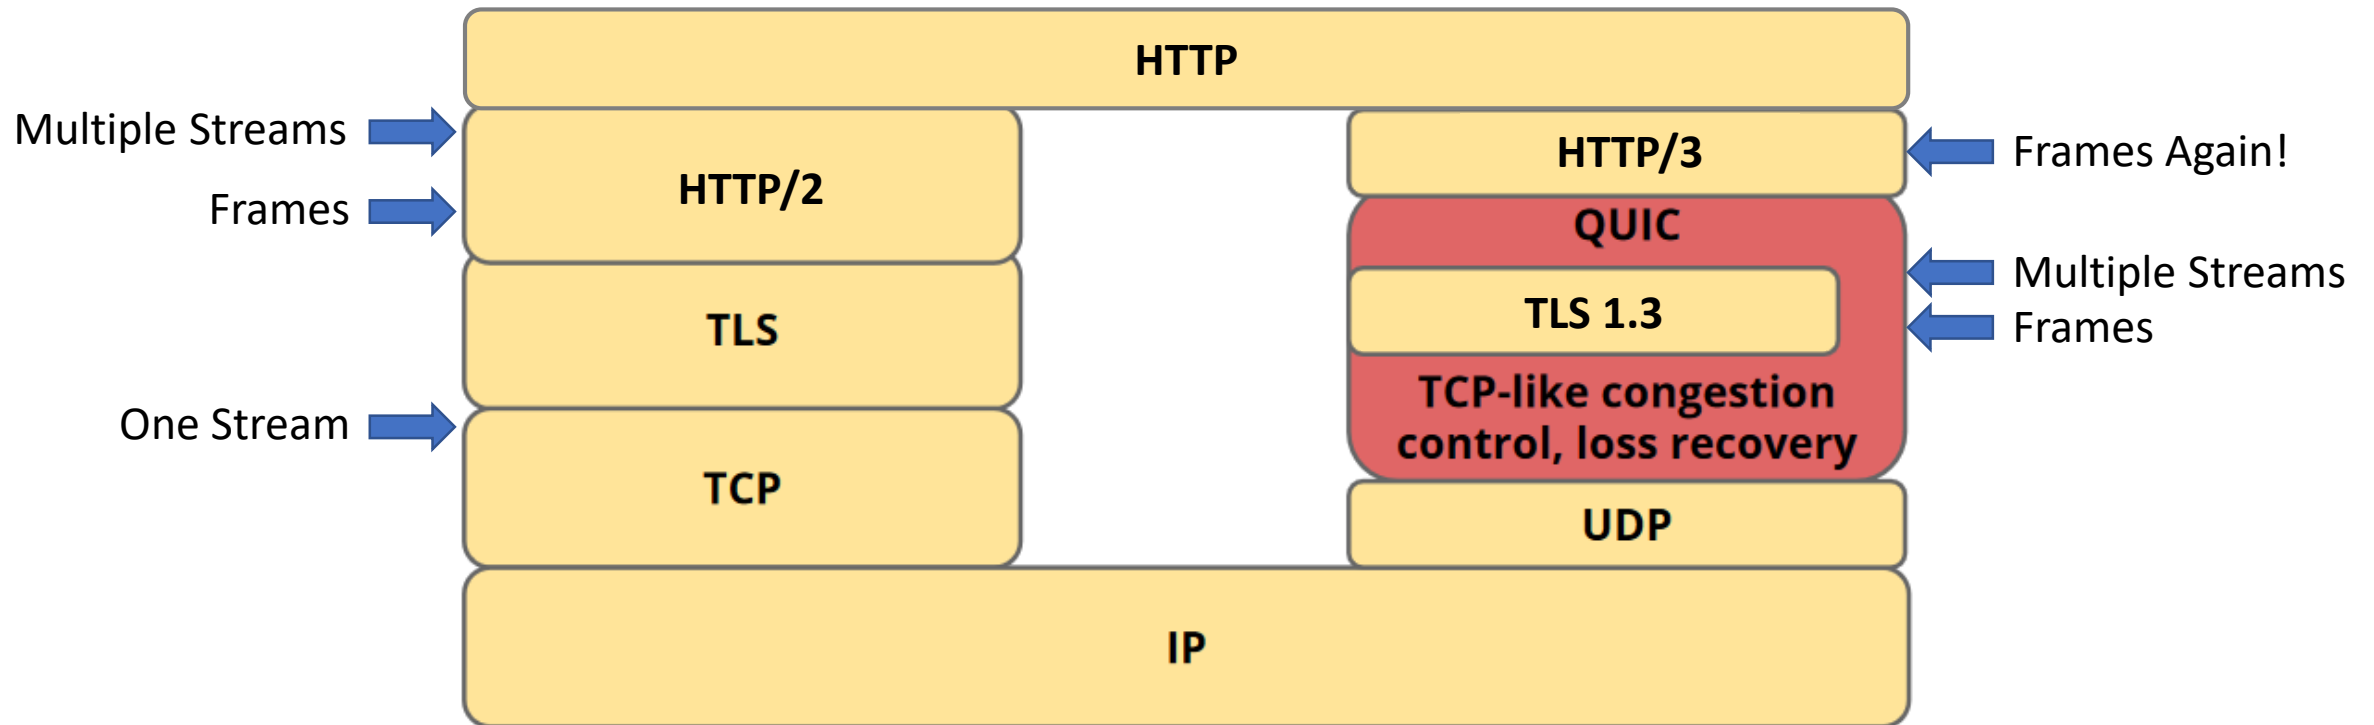
# Multiplexing is a Mixed Bag

- On a low-loss network, HTTP/2 is great:
  - Lower connection setup costs
  - Grow the TCP congestion window larger for better throughput
  - Avoid self-competition with parallel connections
- On a lossy network, HTTP/2 shares vulnerability





# So what's different?



# Ordering is complicated



Within a stream, data is delivered in order



Between streams, no ordering guarantee at all



Inter-stream dependencies risk deadlocks

# Unidirectional Stream Types

Server Bidi	Server Uni
Client Bidi	Client Uni



- HTTP/2 has two varieties; QUIC has four
- Unidirectional streams begin with a type byte
  - If you understand it, keep reading
  - If not, kill the stream (STOP\_SENDING)
- Extensible, like frame types
  - Define frame if data is always a single unit
  - Define stream type if data can develop over time

# HTTP/2 defines ten frame types

DATA

HEADERS

PRIORITY

RST\_  
STREAM

SETTINGS

PUSH\_  
PROMISE

PING

GOAWAY

WINDOW\_  
UPDATE

CONTINUATION

# Every single one is different in HTTP/3

DATA

HEADERS

PRIORITY

CANCEL\_  
PUSH

SETTINGS

PUSH\_  
PROMISE



GOAWAY



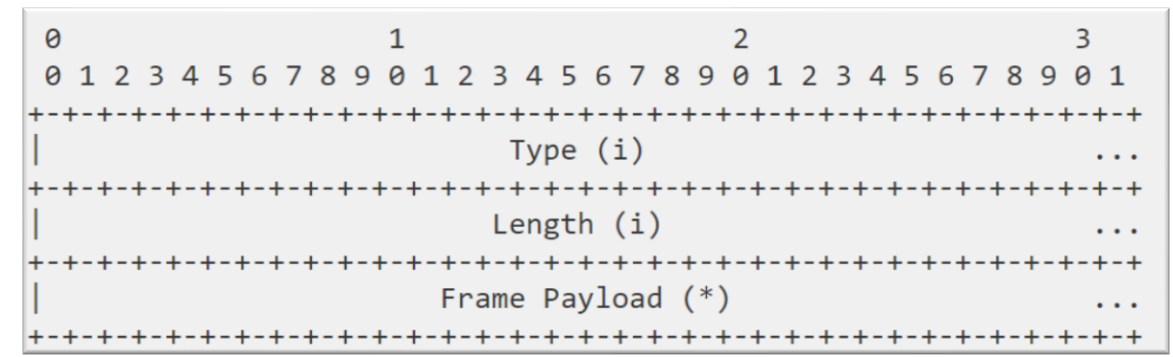
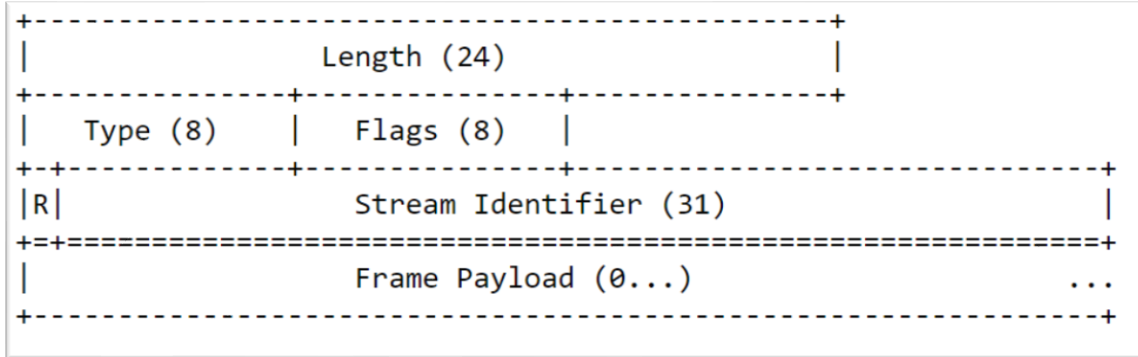
COM...ION

MAX\_  
PUSH\_ID

DUPLICATE  
\_PUSH



# Universal Differences



# SETTINGS

SETTINGS describes what the sender is capable of processing

**Q:** When do changes take effect?

**A:** Send a SETTINGS\_ACK on the control stream describing which streams are currently open, then a SETTING\_ACK on every stream indicating when changes took effect!

# SETTINGS

SETTINGS describes what the sender is capable of processing

**Q:** When do changes take effect?

**A:** ~~Send a SETTINGS\_ACK on the control stream describing which streams are currently open, then a SETTING\_ACK on every stream indicating when changes took effect!~~

No changes to SETTINGS allowed

Assume conservative defaults until you've seen the other side's SETTINGS

HTTP/2 prioritization assumes ordering

HTTP/3 provides the same tree-based concept, with differences:

- Initial priorities are the first frame of the stream
- All priority changes sent on one control stream
- Exclusive prioritization is not possible
- Placeholders and aggressive pruning

Prioritization

# Header Compression

HPACK is effectively a stream of commands:

- Emit this literal value
- Insert this value into the table and use it
- Use the value indexed at #5



# Header Compression

HPACK is effectively a stream of commands:

- Emit this literal value
- Use the value indexed at #5
- Insert this value into the table and use it

A. Frindell, Ed.  
Facebook  
January 29, 2019

# QPACK: Header Compression for HTTP over QUIC

draft-ietf-quic-qpack-latest

## Abstract

This specification defines QPACK, a compression format for efficiently representing HTTP header fields to be used in HTTP/3

# QPACK in 20 seconds

The HEADERS frames are stand-alone:

- Start from table state #28
- Emit this literal value
- Use the value indexed at #5

The table is managed by a stream of commands:

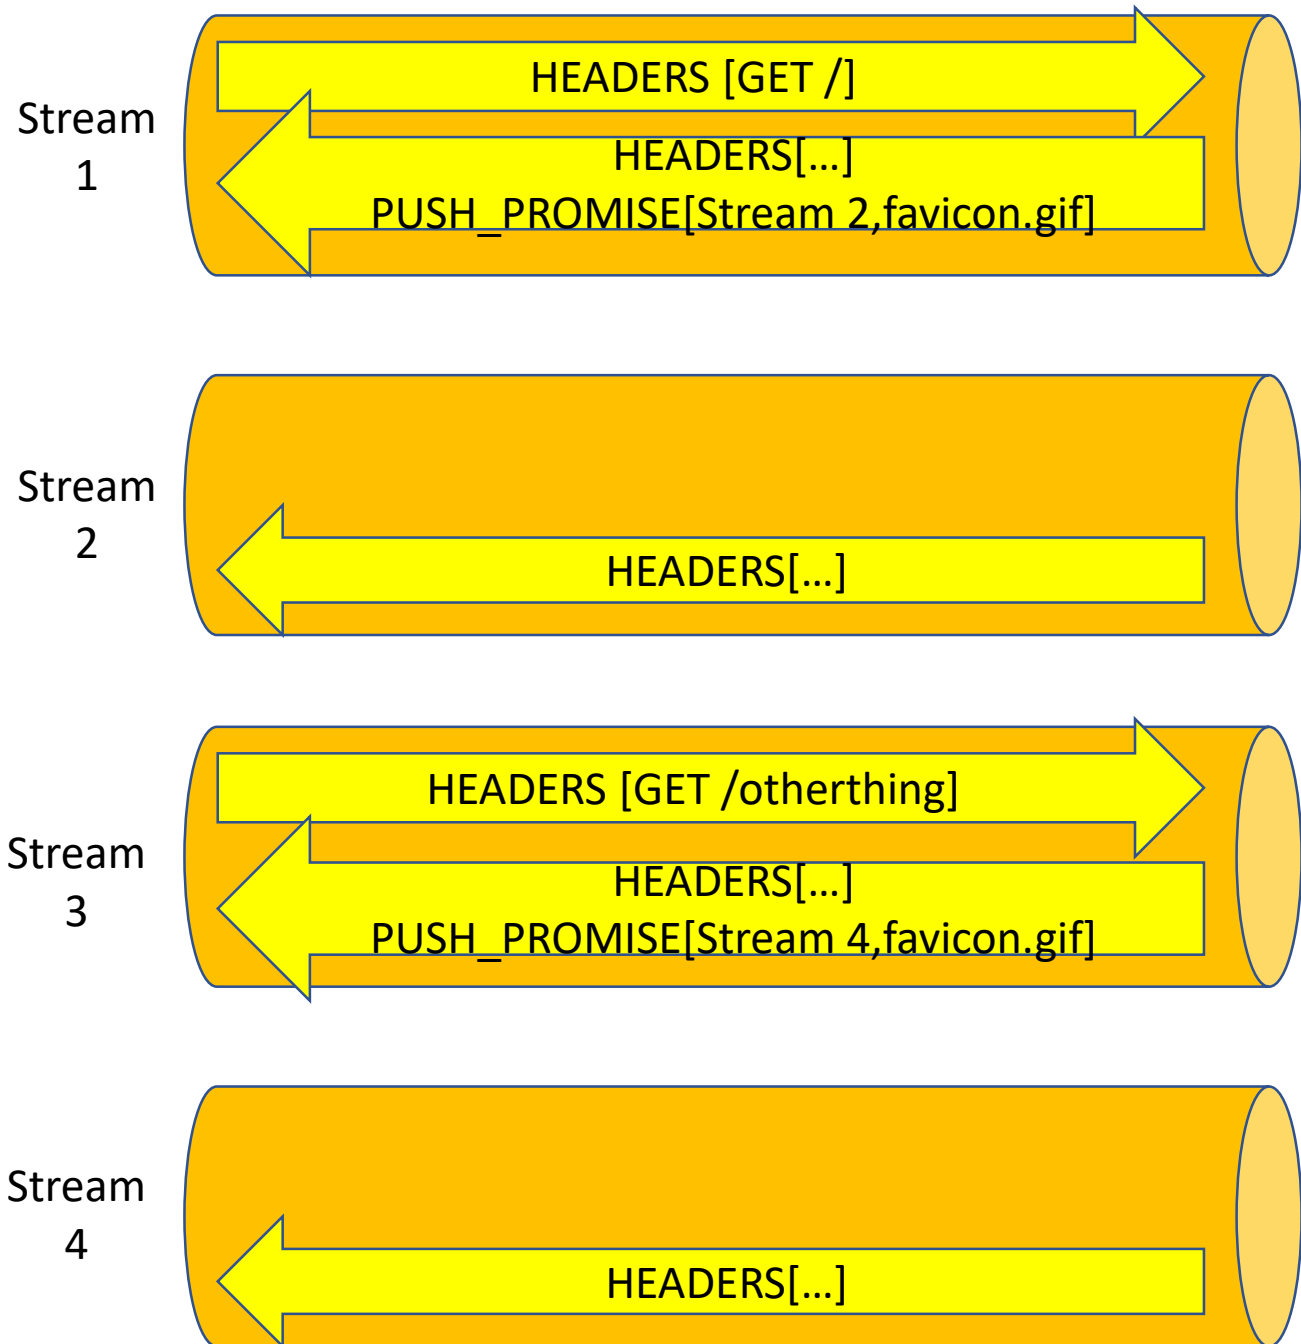
- Insert this value into the table

...and a return stream of status:

- Processed header block on stream #16
- Abandoned reading on stream #12

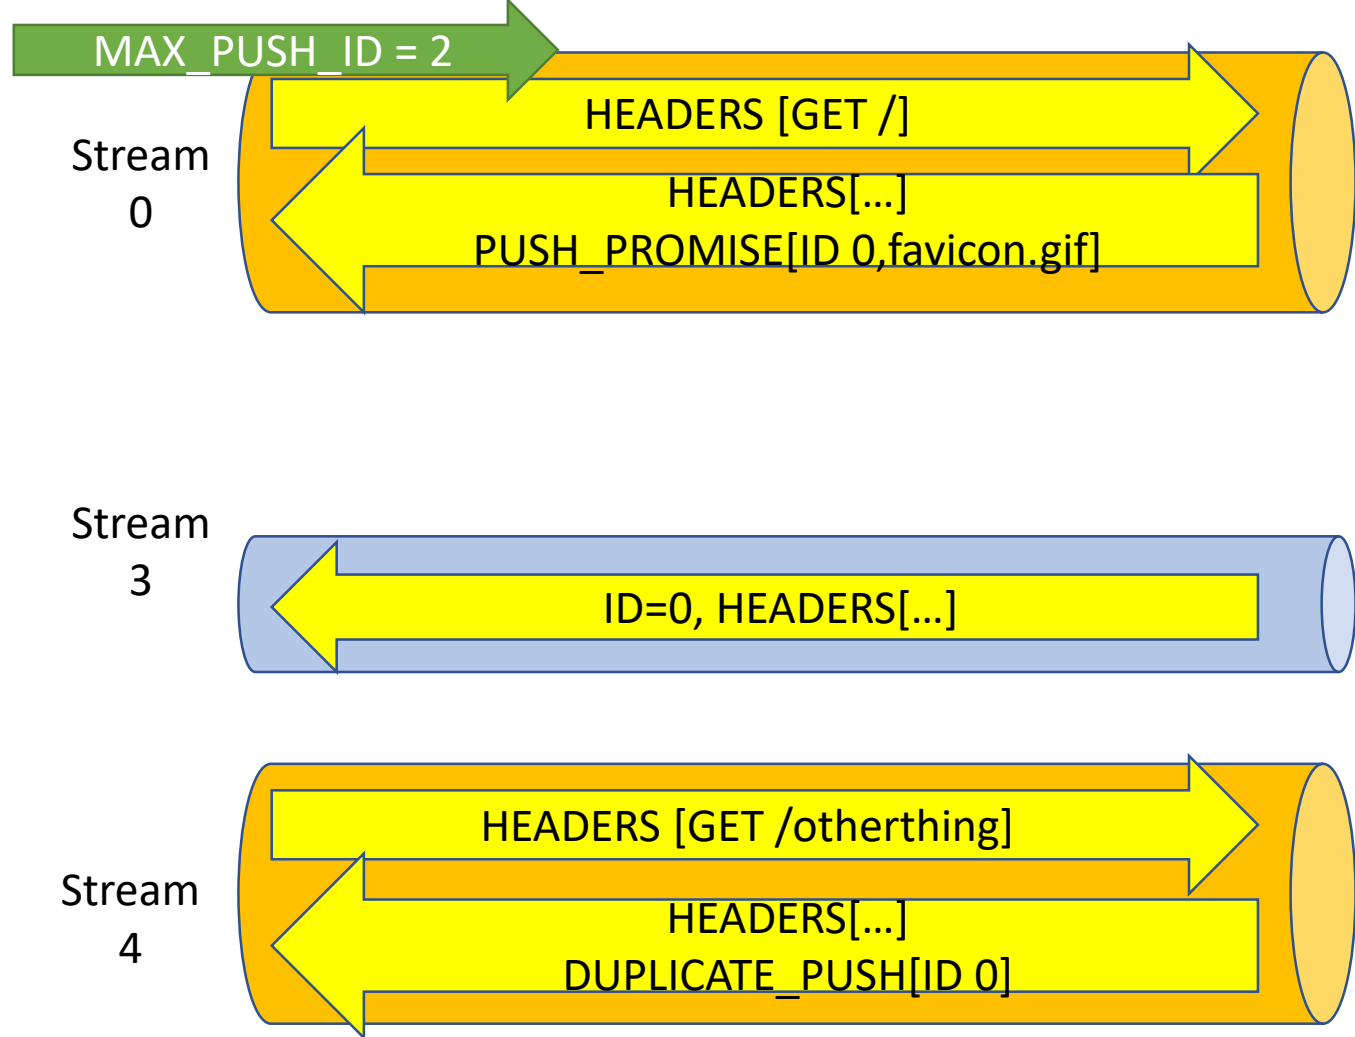
# Server Push HTTP/2

- Push streams are “special” in the stream lifecycle
- Promise identifies stream
- Push controlled by number of streams server is permitted
- Promises and responses match one-to-one



# Server Push HTTP/3

- Push streams are normal unidirectional streams
- ID matches promise to stream
- Push controlled by Push ID limit
- DUPLICATE\_PUSH frame allows many-to-one pushes





# Threats to HTTP/3



QUIC traffic looks an awful lot like DoS traffic



CPU overhead



Ossification inside the box



Necessary middleboxes